

## **Lab Task-4: Weather data analysis for analysing hot and cold days using MapReduce.**

**Submitted by:**

Asish Karthikeya

CSE\_B

### **Code**

#### **Mapper**

```
package weather_data;

import java.io.IOException;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;

public class weather_mapper extends Mapper<LongWritable, Text, Text, IntWritable>
{
    private static final int MISSING = 9999;

    @Override
    public void map(LongWritable key, Text value, Context context) throws IOException,
        InterruptedException
    {
        String line = value.toString();
        String year = line.substring(15, 19);
        int airTemperature;
        if (line.charAt(87) == '+')
        { // parseInt doesn't like leading plus signs
            airTemperature = Integer.parseInt(line.substring(88, 92));
        }
        else
        {
            airTemperature = Integer.parseInt(line.substring(87, 92));
        }
    }
}
```

```
String quality = line.substring(92, 93);  
if (airTemperature != MISSING && quality.matches("[01459]"))  
{  
    context.write(new Text(year), new IntWritable(airTemperature));  
}  
}  
}
```

## **Reducer**

```
package weather_data;  
  
import java.io.IOException;  
  
import org.apache.hadoop.io.IntWritable;  
  
import org.apache.hadoop.io.Text;  
  
import org.apache.hadoop.mapreduce.Reducer;  
  
public class weather_reducer extends Reducer<Text, IntWritable, Text, IntWritable>  
{  
  
    @Override  
  
    public void reduce(Text key, Iterable<IntWritable> values, Context context) throws IOException,  
        InterruptedException  
    {  
  
        int maxVal = Integer.MIN_VALUE;  
  
        for (IntWritable value : values)  
        {  
  
            maxVal = Math.max(maxVal, value.get());  
  
        }  
  
        context.write(key, new IntWritable(maxVal));  
  
    }  
}
```

## **Driver program**

```
package weather_data;  
  
import org.apache.hadoop.fs.Path;  
  
import org.apache.hadoop.io.IntWritable;  
  
import org.apache.hadoop.io.Text;  
  
import org.apache.hadoop.mapreduce.Job;  
  
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;  
  
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
```

```

public class weather_main
{
    public static void main(String[] args) throws Exception
    {
        if (args.length != 2)
        {
            System.err.println("Usage: MaxTemperature <input path> <output path>");
            System.exit(-1);
        }
        @SuppressWarnings("deprecation")
        Job job = new Job();
        job.setJarByClass(weather_main.class);
        job.setJobName("Max temperature");
        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
        job.setMapperClass(weather_mapper.class);
        job.setReducerClass(weather_reducer.class);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);
        System.exit(job.waitForCompletion(true) ? 0 : 1);
    }
}

```

Dataset Used:

## Index of /pub/data/noaa/2022

<a href="#">Name</a>	<a href="#">Last modified</a>	<a href="#">Size</a>	<a href="#">Description</a>
<a href="#">Parent Directory</a>		-	
<a href="#">010010-99999-2022.gz</a>	2023-05-18 18:08	122K	
<a href="#">010014-99999-2022.gz</a>	2023-05-18 18:08	183K	

Input:

## File Preview

/weather\_analysis/input/010010-99999-2022

```
00930100109999920220101004+70939-008669FM-12+001099999V0203581N01391999999999999999-00871-01111101401ADDKA1010M-007
00930100109999920220101004+70939-008669FM-12+001099999V0203531N01271999999999999999-00901-01141101411ADDKA1010M-008
00930100109999920220101004+70939-008669FM-12+001099999V0203501N01051999999999999999-00991-01271101381ADDKA1010M-009
00930100109999920220101004+70939-008669FM-12+001099999V0203511N00851999999999999999-01051-01411101341ADDKA1010M-010
00930100109999920220101004+70939-008669FM-12+001099999V0203491N00791999999999999999-01051-01581101321ADDKA1010M-010
01920100109999920220101004+70939-008669FM-12+001099999V0200021N00681006001C9020000199-01071-01731101321ADDDAY171031AY2
009301001099999202201011004+70939-008669FM-12+001099999V0203431N00891999999999999999-01171-01581101341ADDKA1010M-010
009301001099999202201011004+70939-008669FM-12+001099999V0203441N00601999999999999999-01151-01621101331ADDKA1010M-011
023001001099999202201011004+70939-008669FM-12+001099999V0203401N00821003001C9008000199-01241-01761101251ADDDAY181061AY2
009301001099999202201011004+70939-008669FM-12+001099999V0203401N00741999999999999999-01251-01931101191ADDKA1010M-012
009301001099999202201011004+70939-008669FM-12+001099999V0203391N00751999999999999999-01271-01881101071ADDKA1010M-012
022401001099999202201011004+70939-008669FM-12+001099999V0203411N00721006001C9018000199-01251-01911100991ADDDAY181031AY2
009301001099999202201011004+70939-008669FM-12+001099999V0203561N00641999999999999999-01241-01791100871ADDKA1010M-012
009301001099999202201011004+70939-008669FM-12+001099999V0203481N00761999999999999999-01201-01791100761ADDKA1010M-012
024101001099999202201011004+70939-008669FM-12+001099999V0203531N00671006001C9018000199-01171-01831100671ADDDAA112000231
009301001099999202201011004+70939-008669FM-12+001099999V0203411N00681999999999999999-01171-01771100611ADDKA1010M-011
009301001099999202201012004+70939-008669FM-12+001099999V0203411N00891999999999999999-01101-01451100481ADDKA1010M-011
```

✕ Cancel

Download

## Jar File Path & Location:

### Path:

C:\PBDM\Lab

### Location:

Name	Date modified
New folder	05-09-2023 22:49
hadoop-common-2.5.0	05-09-2023 22:16
hadoop-mapreduce-client-core-2.5.0	05-09-2023 22:16
MatrixMul	05-09-2023 22:18
Weather_analysis	18-09-2023 20:44
Word_Count_Program	04-09-2023 21:07

## Execution:

```
[root@sandbox ~]# hadoop fs -get /weather_analysis/Weather_analysis.jar
get: 'Weather_analysis.jar': File exists
[root@sandbox ~]# hadoop jar Weather_analysis.jar /weather_analysis/input /weather_
23/09/18 10:35:23 INFO impl.TimelineClientImpl: Timeline service address: http://sa
23/09/18 10:35:23 INFO client.RMPProxy: Connecting to ResourceManager at sandbox.hor
23/09/18 10:35:24 INFO client.AHSPProxy: Connecting to Application History server at
23/09/18 10:35:25 WARN mapreduce.JobResourceUploader: Hadoop command-line option pa
oolRunner to remedy this.
23/09/18 10:35:26 INFO input.FileInputFormat: Total input paths to process : 1
23/09/18 10:35:26 INFO lzo.GPLNativeCodeLoader: Loaded native gpl library
23/09/18 10:35:26 INFO lzo.LzoCodec: Successfully loaded & initialized native-lzo 1
23/09/18 10:35:27 INFO mapreduce.JobSubmitter: number of splits:1
23/09/18 10:35:27 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_16945
23/09/18 10:35:28 INFO impl.YarnClientImpl: Submitted application application_16945
23/09/18 10:35:29 INFO mapreduce.Job: The url to track the job: http://sandbox.hort
23/09/18 10:35:29 INFO mapreduce.Job: Running job: job_1694535283099_0010
23/09/18 10:36:00 INFO mapreduce.Job: Job job_1694535283099_0010 running in uber mo
23/09/18 10:36:00 INFO mapreduce.Job: map 0% reduce 0%
23/09/18 10:36:27 INFO mapreduce.Job: map 100% reduce 0%
```

```
23/09/18 10:37:16 INFO mapreduce.Job: map 100% reduce 100%
23/09/18 10:37:18 INFO mapreduce.Job: Job job_1694535283099_0010 failed with state
Job failed as tasks failed. failedMaps:1 failedReduces:0
```

```
23/09/18 10:37:18 INFO mapreduce.Job: Counters: 13
  Job Counters
    Failed map tasks=4
    Killed reduce tasks=1
    Launched map tasks=4
    Other local map tasks=3
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=63769
    Total time spent by all reduces in occupied slots (ms)=0
    Total time spent by all map tasks (ms)=63769
    Total time spent by all reduce tasks (ms)=0
    Total vcore-milliseconds taken by all map tasks=63769
    Total vcore-milliseconds taken by all reduce tasks=0
    Total megabyte-milliseconds taken by all map tasks=15942250
    Total megabyte-milliseconds taken by all reduce tasks=0
[root@sandbox ~]# hadoop jar Weather_analysis.jar /weather_analysis/input /weather_
23/09/18 10:39:51 INFO impl.TimelineClientImpl: Timeline service address: http://sa
23/09/18 10:39:51 INFO client.RMPProxy: Connecting to ResourceManager at sandbox.hor
23/09/18 10:39:52 INFO client.AHSPProxy: Connecting to Application History server at
23/09/18 10:39:53 WARN mapreduce.JobResourceUploader: Hadoop command-line option pa
oolRunner to remedy this.
23/09/18 10:39:54 INFO input.FileInputFormat: Total input paths to process : 1
23/09/18 10:39:54 INFO lzo.GPLNativeCodeLoader: Loaded native gpl library
23/09/18 10:39:54 INFO lzo.LzoCodec: Successfully loaded & initialized native-lzo 1
23/09/18 10:39:56 INFO mapreduce.JobSubmitter: number of splits:1
23/09/18 10:39:57 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_16945
23/09/18 10:39:57 INFO impl.YarnClientImpl: Submitted application application_16945
23/09/18 10:39:58 INFO mapreduce.Job: The url to track the job: http://sandbox.hort
23/09/18 10:39:58 INFO mapreduce.Job: Running job: job_1694535283099_0011
23/09/18 10:40:22 INFO mapreduce.Job: Job job_1694535283099_0011 running in uber mo
23/09/18 10:40:22 INFO mapreduce.Job: map 0% reduce 0%
23/09/18 10:40:44 INFO mapreduce.Job: map 100% reduce 0%
23/09/18 10:41:04 INFO mapreduce.Job: map 100% reduce 100%
23/09/18 10:41:06 INFO mapreduce.Job: Job job_1694535283099_0011 completed successf
ully
23/09/18 10:41:06 INFO mapreduce.Job: Counters: 49
  File System Counters
    FILE: Number of bytes read=44248
```

23/09/18 10:41:06 INFO mapreduce.Job: Counters: 49

File System Counters

FILE: Number of bytes read=44248  
FILE: Number of bytes written=376859  
FILE: Number of read operations=0  
FILE: Number of large read operations=0  
FILE: Number of write operations=0  
HDFS: Number of bytes read=918721  
HDFS: Number of bytes written=8  
HDFS: Number of read operations=6  
HDFS: Number of large read operations=0  
HDFS: Number of write operations=2

Job Counters

Launched map tasks=1  
Launched reduce tasks=1  
Data-local map tasks=1  
Total time spent by all maps in occupied slots (ms)=19604  
Total time spent by all reduces in occupied slots (ms)=16484  
Total time spent by all map tasks (ms)=19604  
Total time spent by all reduce tasks (ms)=16484  
Total vcore-milliseconds taken by all map tasks=19604  
Total vcore-milliseconds taken by all reduce tasks=16484  
Total megabyte-milliseconds taken by all map tasks=4901000  
Total megabyte-milliseconds taken by all reduce tasks=4121000

Map-Reduce Framework

Map input records=4022  
Map output records=4022  
Map output bytes=36198  
Map output materialized bytes=44248  
Input split bytes=141  
Combine input records=0  
Combine output records=0  
Reduce input groups=1  
Reduce shuffle bytes=44248  
Reduce input records=4022  
Reduce output records=1  
Spilled Records=8044  
Shuffled Maps =1  
Failed Shuffles=0  
Merged Map outputs=1  
GC time elapsed (ms)=627  
CPU time spent (ms)=6220

```

Total megabyte milliseconds taken by all reduce tasks=4221000
Map-Reduce Framework
  Map input records=4022
  Map output records=4022
  Map output bytes=36198
  Map output materialized bytes=44248
  Input split bytes=141
  Combine input records=0
  Combine output records=0
  Reduce input groups=1
  Reduce shuffle bytes=44248
  Reduce input records=4022
  Reduce output records=1
  Spilled Records=8044
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=627
  CPU time spent (ms)=6220
  Physical memory (bytes) snapshot=297906176
  Virtual memory (bytes) snapshot=3891314688
  Total committed heap usage (bytes)=139984896
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=918580
File Output Format Counters
  Bytes Written=8
[root@sandbox ~]# hadoop fs -cat /weather_analysis/output/part-r-00000
2022      92
[root@sandbox ~]#

```

## Output:

### File Preview

/weather\_analysis/output/part-r-00000

```
2022      92
```